

Consensus

Thierry Sans

The consensus problem

In lecture 2, we said

"one node is selected to create the next block"

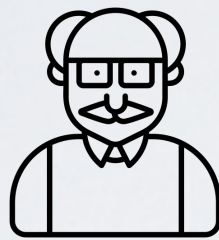
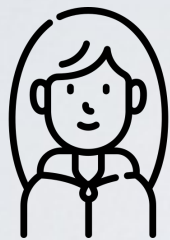
But how nodes agree on something in which some nodes in the P2P network might fail or not be honest?

Outline

- Byzantine Agreement Problem and the Sybil attack
- Proof of Work (a.k.a Nakamoto Consensus)
- Proof of Stake
- Other consensus protocols

The Byzantine Agreement Problem and The Sybil Attack

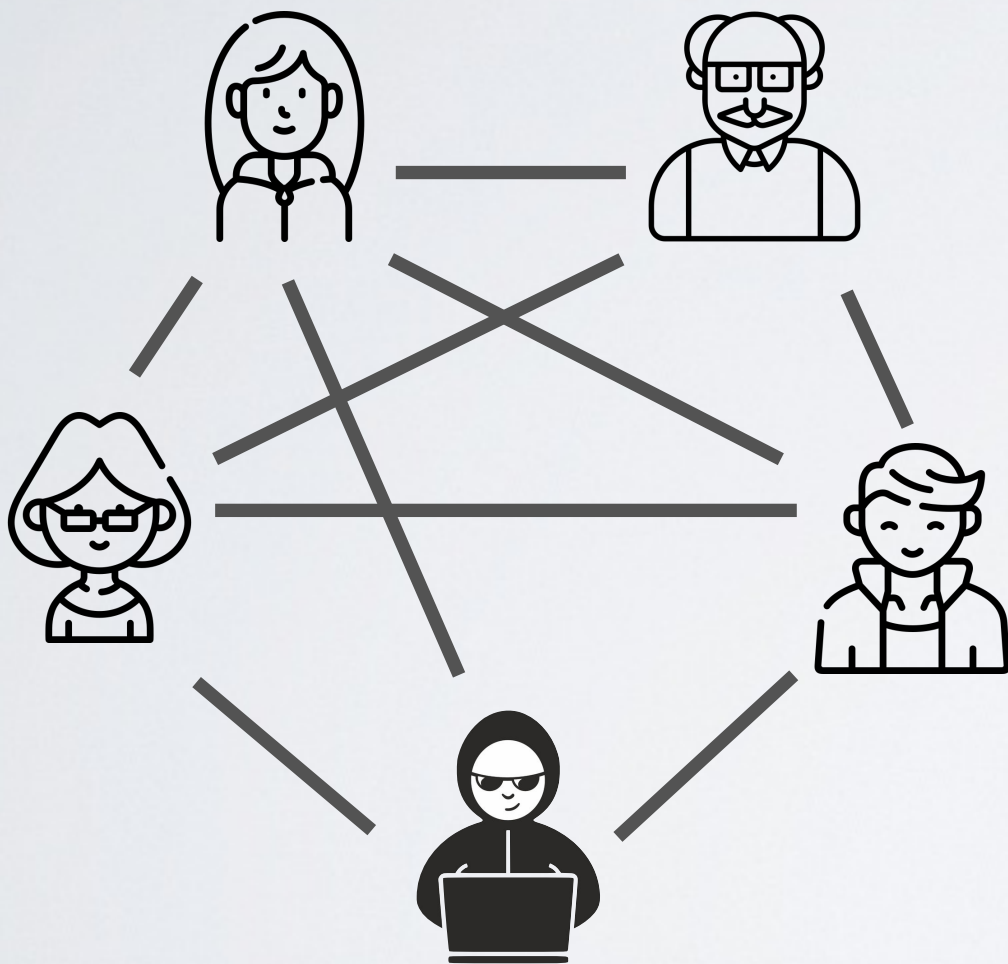
The byzantine agreement problem



Generals attacking a fortress must decide whether to attack or retreat

- ✓ They must reach a common decision and make a coordinated action (consensus)
- Some participants might default or sabotage the consensus process (a.k.a Byzantine failure)

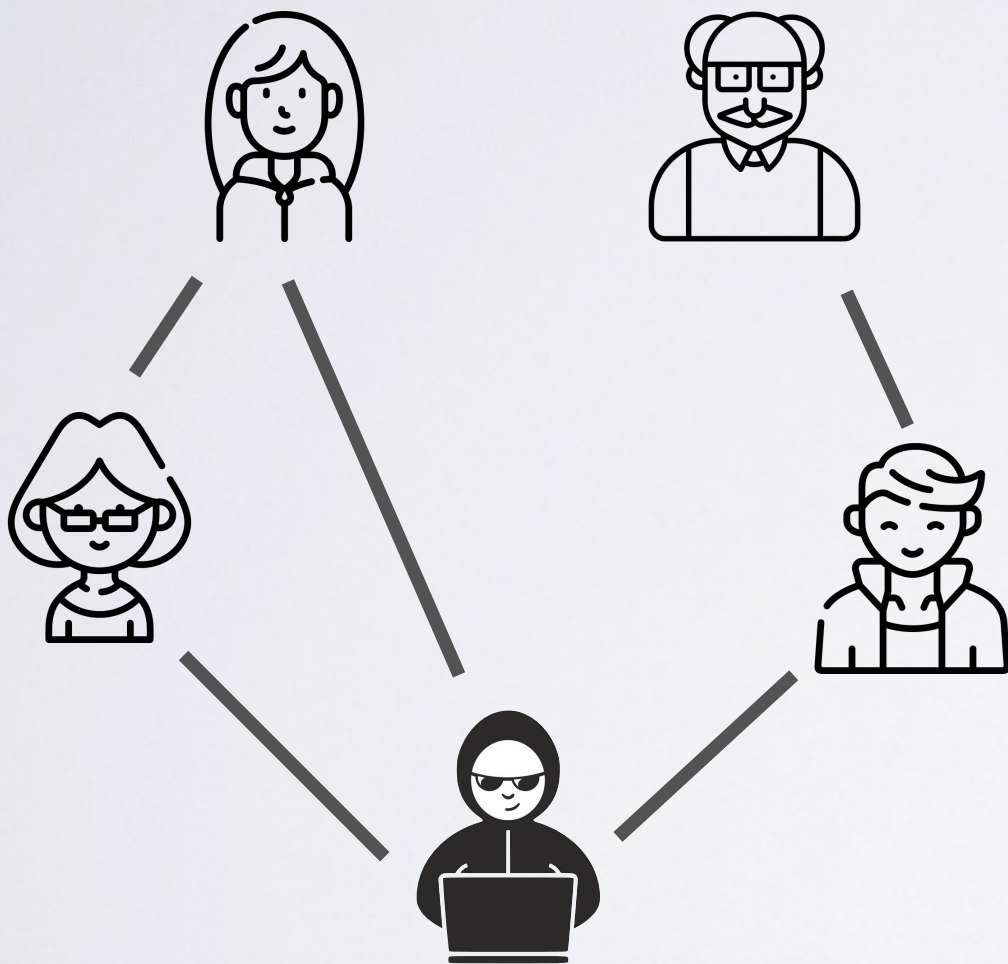
A fully connected network



Each general sends their vote to all others individually

- (failure) some votes might not reach their destination
- (sabotage) a general might send different vote to different people

A P2P network



Each general signs and forwards their vote to peers that we relay it (flooding algorithm)

- (failure) some votes might not reach their destination
- (sabotage) a general might send different vote to different people
- (sabotage) a general might not forward certain vote

Requirements for Byzantine Fault Tolerance (BFT)

Given a system of n nodes, t of which are dishonest. When a node A broadcasts the value x , the other nodes are allowed to discuss with each other and verify the consistency of A 's broadcast, and eventually settle on a common value y .

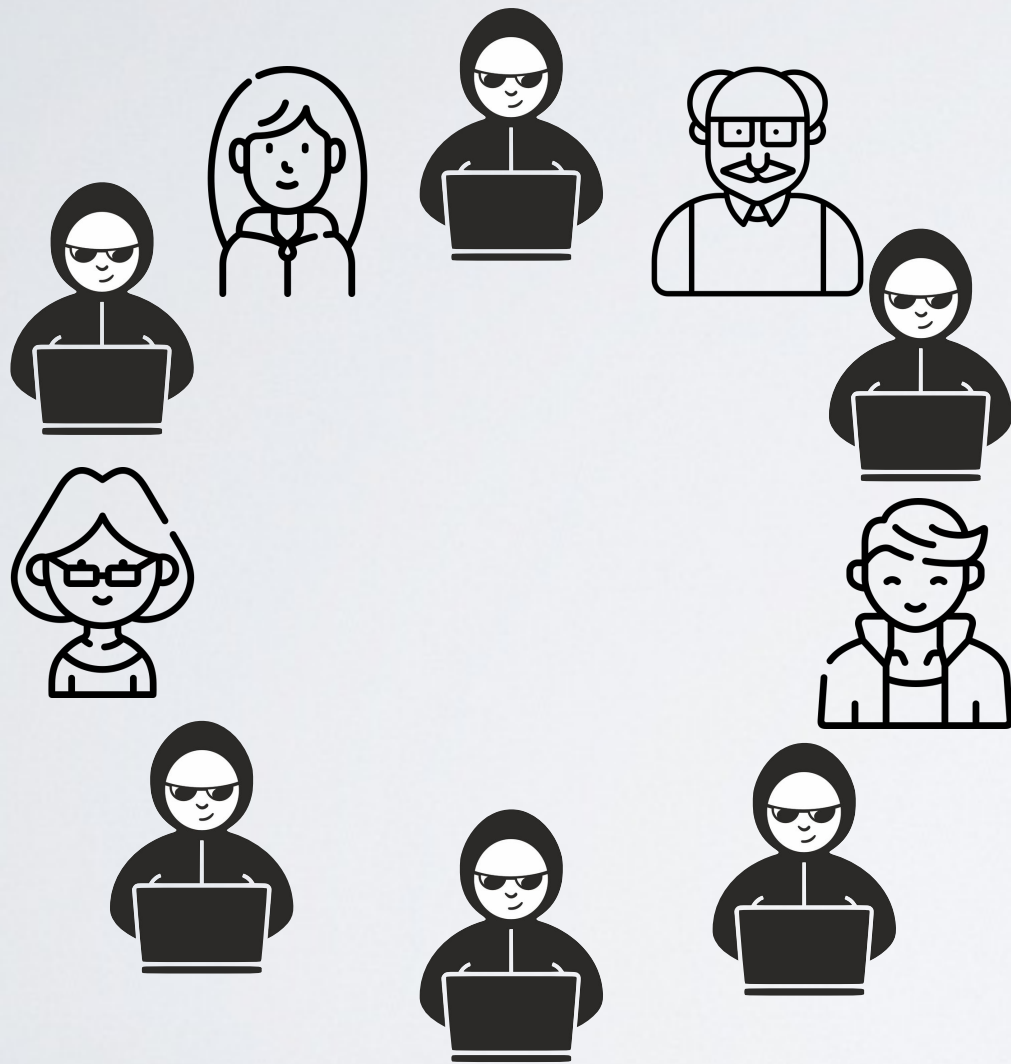
The system is said to resist Byzantine faults if either

- A is honest and all honest nodes agree that A says x , or
- A fails or is dishonest and all honest node agree that A says Y

Security properties:

- **Consistency** - Honest nodes do not contradict
- **Liveness** - Progress is made

Another problem - Sybil Attack



If a reputation system, an attacker can create and operate multiple nodes/identities to

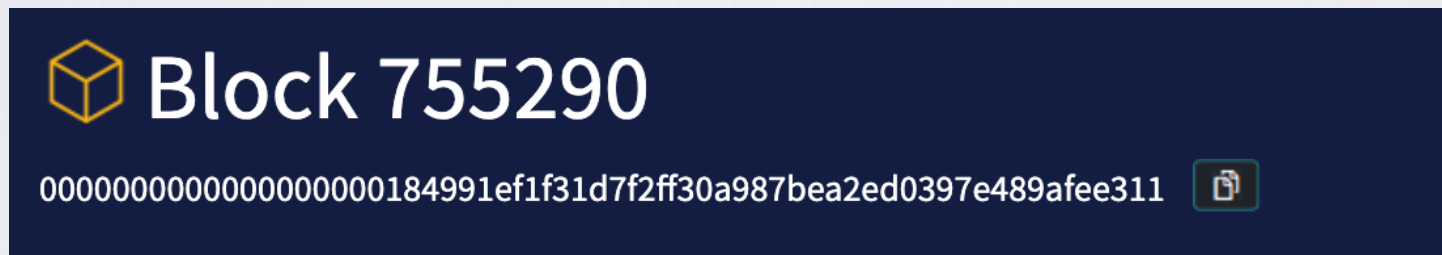
- Carry out a 51% attack to control the consensus outcome
- Block messages from honest nodes (P2P network)

Proof of Work (a.k.a Nakamoto Consensus)

The Mining Challenge



➔ Find a nonce such that $H(\text{block})$ has a certain number of leading 0s



The number of zeros is determined from the target

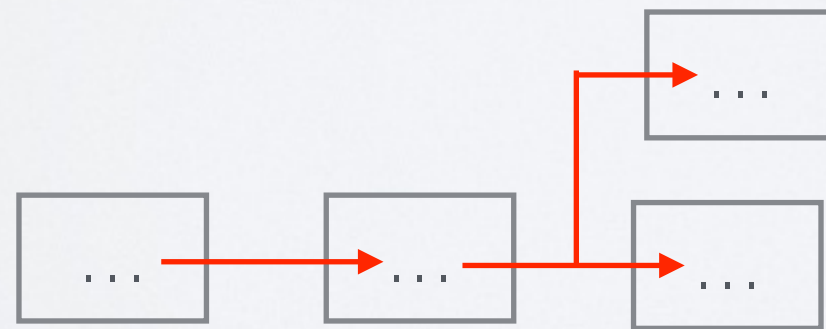
https://en.bitcoinwiki.org/wiki/Difficulty_in_Mining#Bitcoin_mining_difficulty

Target is adjusted every 2016 blocks (every 2 weeks) to mine a block every 10 minutes in average

- ✓ Finding the right nonce that outputs the right hash is hard (but not impossible) to compute
- ✓ Verifying the nonce and the block hash is easy to verify by other nodes

Good but ...

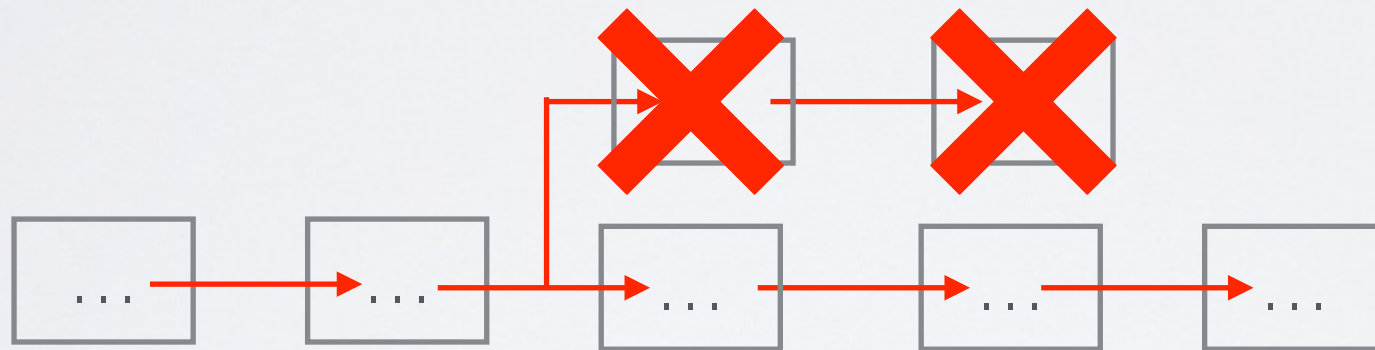
- ✓ Prevents the sybil attack
Probability of winning is relative to the computation power
- Does not fully solve the Byzantine Agreement Problem (yet)
The mining challenge does prevent **two valid** blocks be found and broadcasted at relatively same time



The Longest-chain Wins Protocol

Simple rule

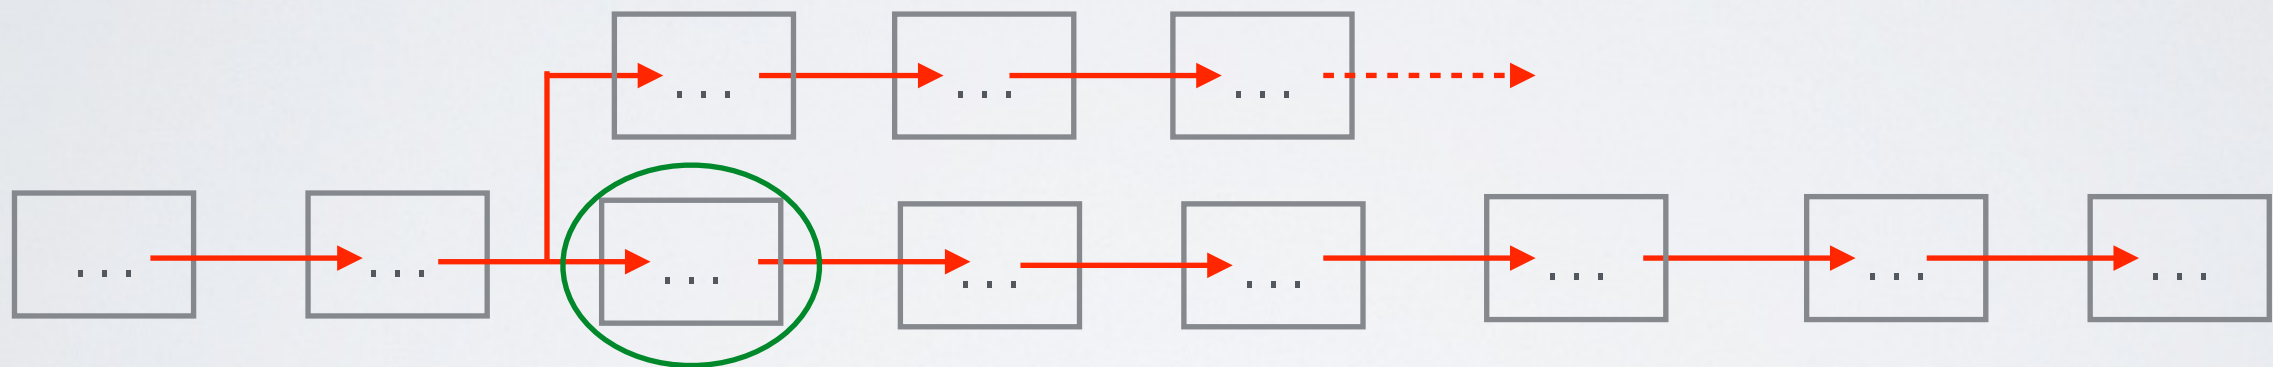
A miner must mine a block on the longest chain known



➔ **Consequence** : one should wait at least 6 confirmation blocks to consider a transaction as fully confirmed

Double-spending attack

Can Mallory perform a double-spending attack by crafting a chain that exceeds 6 blocks previously confirmed?



- ✓ Binomial Random Walk (see Bitcoin paper)
Impossible to catch up with the longest chain unless Mallory controls more than 51% of the blockchain computing power

Criticism of Proof of Work

- All miners spend considerable amount of energy to eventually have only one of them being rewarded

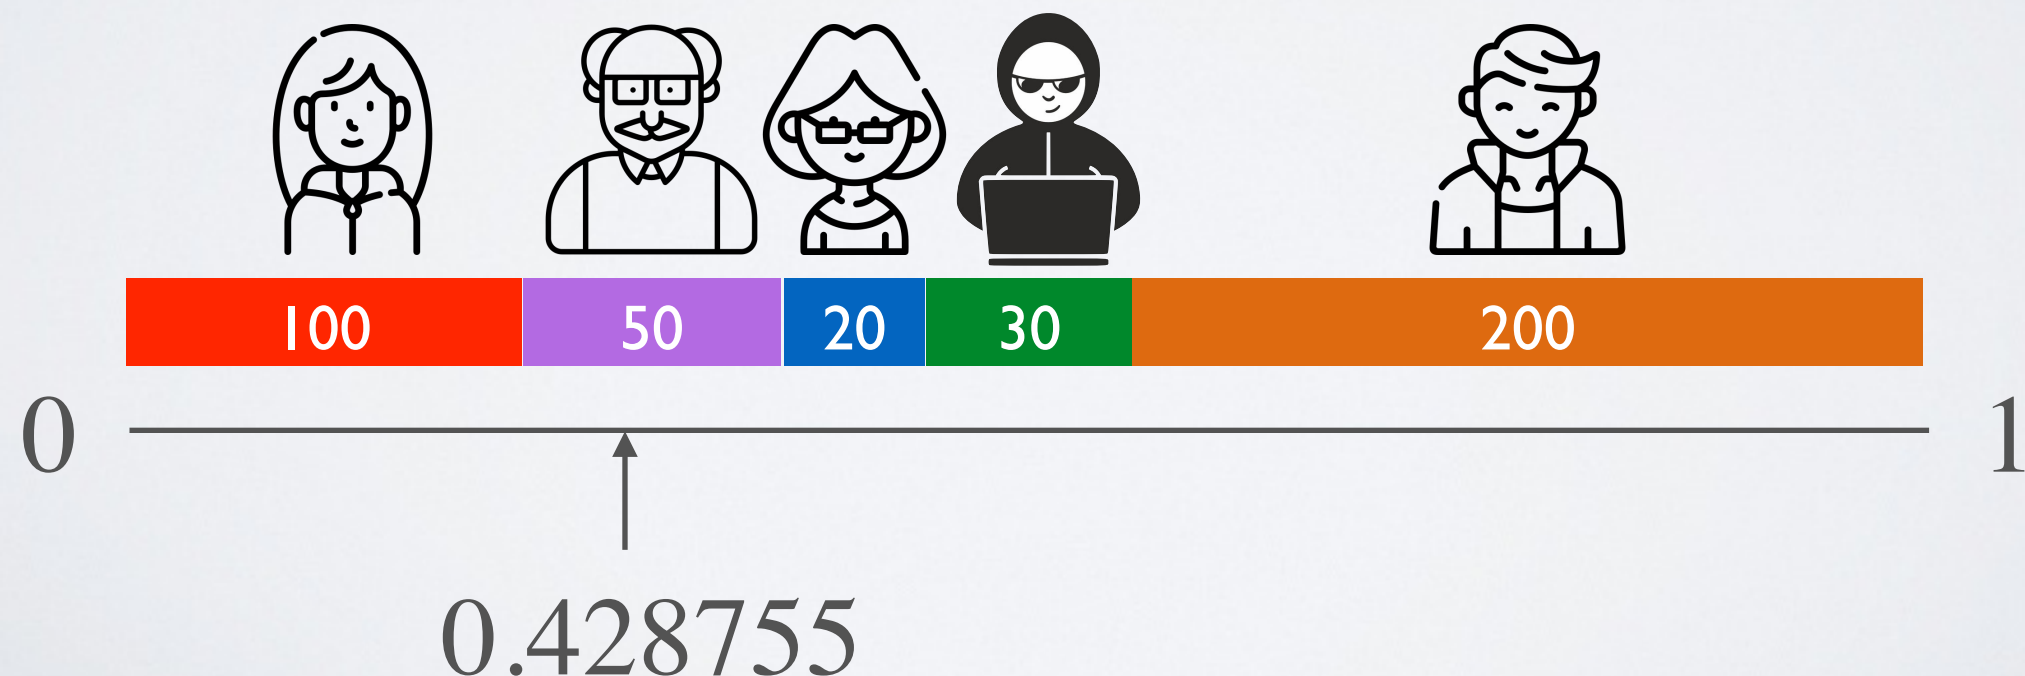
Proof of Stake

A more energy efficient approach


- ✓ Solving the Byzantine Agreement Problem by electing the node that will mine the block
- ✓ Preventing a sybil attack by making the probability of being elected relative to the monetary power (staking)

Election process

Electing one of the node to mine the next block requires **a beacon** : an ideal service that regularly publishes random value which no party can predict or manipulate



Collect Random Approach



Step 0 - Each node i generate a random number r_i

Step 1 - Each node commits its random number by broadcasting $H(r_i)$

Step 2 - Each node broadcasts its random r_i and verify all committed hashes others

Step 3 - Each node calculate the beacon value
$$\text{beacon} = H(r_1 \parallel r_2 \dots \parallel r_n)$$

Step 4 - The elected node mine the next block and broadcast it to the network

Step 5 - The other nodes check the block
If the block is invalid, the elected node lose its stake (slashing)

Beyond Pow and Pos

Other consensus

- Proof of Stake variant based on pBFT - Practical Byzantine Fault Tolerant (Cosmos Tendermint)
- Proof of Authority ~ private network (Facebook Libra)
- Proof of History (Avalanche)
- DGA - Directed Acyclic Graph (IOTA)